

AI Literacy Cheat Sheet

Everything from Module 1 on one page.
Keep it in a tab. You'll know it stuck the day you stop opening it.

5 THE IDEAS EVERYTHING RESTS ON

1 · An LLM predicts the next chunk of text

Trained to produce language that *sounds* right, one token at a time. Fluent and fast, not a database of truth.

PM lever: treat every output as a draft to verify, never a fact to trust.

2 · Tokens & the context window

A token is ~4 characters, the unit you're billed for in and out. The context window is the model's finite short-term memory; overflow silently drops the oldest text.

PM lever: verbosity is a line item; long chats "lose the plot" for a reason.

3 · Hallucination is baked in

Confident, fluent, specifically wrong output, in the gap between "sounds right" and "is right." The model can't tell when it's guessing.

PM lever: design for verification, not perfection. The next release won't "fix" it.

4 · Strengths vs. limits

Great at language & pattern tasks with room to review. Struggles with exact math, recent facts, private data, long multi-step logic, and perfect repeatability.

PM lever: most limits are signals for what to build *around* the model.

5 · Your job is the scaffolding

Connect a calculator, live search, or your own docs (RAG) and the limits fall away. Building an AI feature is mostly the scaffolding, not the model.

PM lever: spec the guardrails, grounding, and human checks, not just the prompt.

9 THE WORDS YOU NEED TO OWN

Prompt	Your instruction + context. Bad output is usually a prompt problem, not a broken model.
System prompt	Hidden rules set before the user types. Building an AI feature is largely designing this.
Token	~4 characters of text. The unit you pay for, in and out.
Context window	The model's finite short-term memory. The hard ceiling on what it can consider at once.
Fine-tuning	Training a model further on your examples. Expensive, data-hungry, goes stale, rarely step one.
RAG retrieval-augmented generation	Fetch real info first, then answer from it. The #1 way to ground answers and cut hallucination.
Agent	AI that takes multiple actions toward a goal. Powerful, and it raises the safety stakes, it acts.
Multimodal	Handles images, audio, video, not just text. Opens product doors text-only never could.
Temperature	The creativity dial: low = consistent, high = varied. Legal summariser low; brainstorm tool high.

PHRASES THAT SHOULD MAKE YOU PAUSE

"The model's just broken." Usually a prompt problem, not the model.

"Let's fine-tune it." As step one? Ask what they tried first, try RAG.

"The next version fixes hallucinations." It won't. Design for verification.

"Just put everything in the prompt." The window is finite, and you pay per token.

TAKE THIS TO ANY MEETING

The two-question test

QUESTION 1

Is the core a language or pattern task?

QUESTION 2

Can we tolerate the occasional error? (low stakes, or grounded + a human checks)

Both yes?

Green light.

Anything else: add grounding + review, or pick another problem.