

The AI PRD

The parts a normal PRD misses. Page 1 is yours to fill in. Page 2 is the same template worked through on the scheduling assistant.

How to use this. This is the delta, not a whole PRD — keep your normal problem, users, scope and rollout sections and add these. Six new, two rewritten. If you have no AI feature at work yet, fill it in for the scheduling assistant this module has used throughout: the user types **“I need thirty minutes with these four people this week”** and the product proposes slots, explains its choices and drafts the invite. Page two is that one, worked through.

SIX SECTIONS A NORMAL PRD DOES NOT HAVE

1. WHAT GOOD LOOKS LIKE

Describe a good answer the way a person would recognise one, then say who or what decides. No adjectives you cannot test.

2. FAILURE MODES, AND THE BEHAVIOUR FOR EACH

One row per failure. Retrieval returns nothing. The action half-succeeds. The model is confidently wrong. Beside each: what the product does, and what the user sees.

3. GROUNDING

What is it allowed to know, and from where? Which sources, how fresh, and what happens when the source is missing or contradicts itself.

4. THE COST AND LATENCY BUDGET

Cost per request, cost at your forecast volume, and the latency the user feels. Name the levers you have left: cache, batch, route.

5. THE HUMAN IN THE LOOP

Plot it on the square: how bad, and can it be undone. Severe and irreversible gets a person before the action, not a warning after it.

6. WHAT GETS WATCHED AFTER LAUNCH

Which numbers, who looks at them, how often, and what number makes you turn it off.

TWO SECTIONS THAT CHANGE

ACCEPTANCE CRITERIA — NOT PASS/FAIL ANY MORE

Real examples, a threshold, and a named judge. “N of M examples judged acceptable by X.”

THE SUCCESS METRIC — CHECK IT CANNOT BE GAMED BY FAILING

If the model can move this number by being wrong, it is the wrong number. Pair it with an outcome metric.

The AI PRD, worked

The scheduling assistant: the user types “I need thirty minutes with these four people this week” and the product proposes slots, explains its choices and drafts the invite.

Worked example. The scheduling assistant, specced. Your own feature replaces it anywhere. The point of reading this one is the **shape** of the answers — concrete, numbered, and decided in advance.

1. WHAT GOOD LOOKS LIKE

A good answer proposes **up to three slots** that are genuinely free for all four people in their own time zones, in the next five working days, and says **in one line why** it picked each one (“only slot where nobody has back-to-back meetings”). **Judged by:** the calendars themselves — a proposed slot is either free for all four or it is not.

2. FAILURE MODES, AND THE BEHAVIOUR FOR EACH

Calendar fetch returns nothing for one person → do not guess; say which person we could not read and offer to send them a poll. **The invite half-writes** (two of four accepted, the write fails) → roll back all four, tell the user, never leave a partial meeting. **Confidently wrong slot** (a stale calendar) → every proposed slot carries the timestamp it was read, and the invite re-checks at send.

3. GROUNDING

Allowed to read: the four attendees’ free/busy, their working hours and time zones, meeting-room availability. **Not allowed to read:** event titles, descriptions or attachments — they are untrusted text a stranger can write into. Freshness: re-read at send, not just at propose.

4. THE COST AND LATENCY BUDGET

~1,200 tokens in, ~300 out per request → **\$0.0006 a request** on the cheap tier. 40 requests per customer per month → **2.4¢**. At 2,000 customers, **about \$48 a month**. On the frontier tier the same feature is **\$2,160 a month**. Latency budget: **first slot on screen in under 3 seconds**. Levers left: cache the system prompt, route the easy single-attendee case to the cheap model, batch nothing (the user is waiting).

5. THE HUMAN IN THE LOOP

Proposing slots is **minor and reversible** → no human. **Writing the invite is irreversible** (it lands in four calendars and four inboxes) → **the user presses send**. The model drafts; it never dispatches.

6. WHAT GETS WATCHED AFTER LAUNCH

Weekly: share of proposals accepted without editing; share where a proposed slot turned out to be busy (the one that must stay near zero); p95 time to first slot; cost per customer. Owner: the PM, in the Monday review. **Turn-it-off number:** busy-slot rate above 2%.

ACCEPTANCE CRITERIA — NOT PASS/FAIL

40 real scheduling requests drawn from last quarter’s support and sales calendars. **At least 36 of 40** produce a slot that is genuinely free for all four attendees. **Judge: the calendars**, checked automatically, not a reviewer’s opinion. The 4 that fail must fail loudly, in one of the ways written in section 2.

THE SUCCESS METRIC — AND WHY IT IS NOT THE OBVIOUS ONE

Not “share of meetings scheduled without a human” — that goes up when the feature fails and the user gives up on it. **Use: meetings scheduled that were not later moved or cancelled within 24 hours**, paired with time-to-scheduled. The model cannot move that one by being wrong.

BEFORE YOU CALL IT FINISHED

Read it as your engineer (what does it do when it goes wrong?), then as finance (what does this cost me if it works?), then as risk (who finds out when it is wrong?). If a section does not answer one of those in their own terms, it is not done. The fourth question — **is it actually any good?** — is the one this template cannot answer for you. That is Module 7.