

The Eval Rubric

Examples, a bar and a number, on one page. Page 1 is yours to fill in, for any AI feature. Page 2 is the same template worked through on the scheduling assistant.

How to use this. It slots into the **What good looks like** section of your AI PRD (6.6), and turns its acceptance criteria into a rate. **Write the rows after reading your traces, not before:** read 30, note each one, name the failures, count them (7.4), and turn the top ones into rows. Each row: **one question** · **binary** · **a pass and a fail example** · **your product's words**. Tick who can grade it: code first, a judge for what is left. Page two is the scheduling assistant, worked through.

THE RUBRIC — FIVE ROWS, EACH ONE A PRODUCT DECISION

FEATURE _____ VERSION ____ DATED _____

ROW 1

PASS EXAMPLE · PASTE A REAL OUTPUT, VERBATIM

FAIL EXAMPLE · PASTE A REAL OUTPUT, VERBATIM

CODE

JUDGE

PERSON

ROW 2

PASS EXAMPLE · PASTE A REAL OUTPUT, VERBATIM

FAIL EXAMPLE · PASTE A REAL OUTPUT, VERBATIM

CODE

JUDGE

PERSON

ROW 3

PASS EXAMPLE · PASTE A REAL OUTPUT, VERBATIM

FAIL EXAMPLE · PASTE A REAL OUTPUT, VERBATIM

CODE

JUDGE

PERSON

ROW 4

PASS EXAMPLE · PASTE A REAL OUTPUT, VERBATIM

FAIL EXAMPLE · PASTE A REAL OUTPUT, VERBATIM

CODE

JUDGE

PERSON

ROW 5

PASS EXAMPLE · PASTE A REAL OUTPUT, VERBATIM

FAIL EXAMPLE · PASTE A REAL OUTPUT, VERBATIM

CODE

JUDGE

PERSON

☐ **Two-grader test:** ten outputs, graded alone by me and by _____ · agreed on ____ of 10 · rows we split on: _____ below 9 of 10, sharpen and re-run

THE GOLDEN SET — ONE ROW PER EXAMPLE · START WITH 30, WORK TOWARD ~100

THE INPUT (verbatim + state)	WHAT GOOD LOOKS LIKE (a condition)	WHY IT IS HERE (the failure it guards)	SOURCE
			REAL · FEARED · SYNTH
			REAL · FEARED · SYNTH
			REAL · FEARED · SYNTH

THE NUMBER — SET IT BEFORE THE FIRST RUN, NOT AFTER

THE THRESHOLD

Ship if ____ of ____ golden cases pass all rows. And what blocks regardless?

WHY THAT NUMBER

First run? Why not 100? What does too low let through?

WHO OWNS IT

One name. Judge last checked against 50 hand-graded: passes ____ / fails ____

SEAT IT IN THE SPEC — IT REWRITES TWO NEIGHBOURS

ACCEPTANCE CRITERIA, AS A RATE

Not "it works". N of M golden cases pass all rows, rubric version and date.

THE QUALITY NUMBER BESIDE YOUR USAGE METRIC

So nobody celebrates adoption of something quietly getting worse.

The Eval Rubric, worked

The scheduling assistant: the user types “I need thirty minutes with these four people this week” and the product proposes slots, explains its choices and drafts the invite.

Worked example. Attendees Priya (London), Jonas (Berlin), Maria (New York) and Tom (Lisbon), working 09:00–17:30 local unless their calendar says otherwise. The rows are the five from Lesson 7.3, word for word. Three are **lookups a plain check settles for free**; two need **a judge**. Your own feature replaces all of it.

THE RUBRIC — SCHEDULING ASSISTANT

VERSION 2 · DATED 12 SEP 2026 · OWNER: THE SCHEDULING PM

ROW 1	Every proposed slot falls inside all four attendees' working hours.	CODE	JUDGE	PERSON
PASS	Tue 14:30 London · 15:30 Berlin, 09:30 New York, 14:30 Lisbon. Inside all four.	FAIL	Wed 16:00 London · 17:00 Berlin, after Jonas's working hours end at 16:30.	
ROW 2	No slot lands before eight or after six in that attendee's own local time.	CODE	JUDGE	PERSON
PASS	Thu 14:00 London · 09:00 in New York, the earliest local time of the four.	FAIL	Thu 12:00 London · 07:00 in New York . The eleven tickets in 7.5.	
ROW 3	The explanation names at least one specific constraint it worked around.	CODE	JUDGE	PERSON
PASS	"Tuesday 14:30 London sits inside the ninety minutes when Maria's day in New York and Jonas's in Berlin overlap."	FAIL	"This works for everyone." Names nothing it worked around.	
ROW 4	If any calendar could not be read, the answer says so explicitly.	CODE	JUDGE	PERSON
PASS	"I could not read Tom's calendar, so these may clash for him. Check with him before sending."	FAIL	Three slots, no mention of Tom. His fetch came back empty in 9 ms; the answer never says so.	
ROW 5	It proposes between two and four options, never one and never nine.	CODE	JUDGE	PERSON
PASS	Three slots, each with its one-line reason.	FAIL	Nine slots in a list. The user now does the choosing the feature exists to do.	

● **Two-grader test:** the PM and the design lead, ten outputs graded alone · first try **6 of 10** on “is the explanation helpful” → after two rewrites, row 3 **10 of 10**

THE GOLDEN SET — THREE OF THE HUNDRED, ONE FROM EACH SOURCE

THE INPUT	WHAT GOOD LOOKS LIKE	WHY IT IS HERE	SOURCE
"30 min with Priya, Jonas, Maria and Tom this week." Tom's calendar returns a permission error.	Any slot inside all four working windows, and the answer says Tom's calendar could not be read.	The 7.4 incident: an empty fetch in 9 ms, a proposed clash, "the model hallucinated".	PRODUCTION FAILURE
"Half an hour with Jonas and Maria tomorrow." Berlin and New York: overlap 15:00–16:30 Berlin only.	Every slot ends by 16:30 Berlin and starts after 15:00. Reason names the time zone gap.	Narrow overlaps are where the 12 "outside working hours" failures cluster.	REAL TRAFFIC, THE ODD ONE
"Thursday the 25th with Priya and Tom." The 25th is a Friday.	Does not silently pick a day. Asks which, or proposes Thursday the 24th and says it assumed that.	One of eight generated ways to write a date wrong. A silent guess books the wrong day.	SYNTHETIC FILL

THE NUMBER

THE THRESHOLD 80 of 100 golden cases pass all five rows Blocks regardless: any case that passed last release and fails now, until the owner has read it.	WHY THAT NUMBER First run came back 73 : healthy, and proof the set is hard. 80 sits above it, so the prompt fix has to land before anything ships. Not 100 : a set everything passes is too easy to warn you. The re-run scored 81 with two locked-calendar regressions , so the block rule, not the headline, held the release.	WHO OWNS IT · THE JUDGE Rows 1, 2, 5 by code; rows 3, 4 by a judge. Judge checked on 50 the PM graded by hand: 24 of 25 passes agreed, 21 of 25 fails agreed (illustrative). 25 a month re-graded by the PM, for good.
---	--	--

SEAT IT IN THE SPEC — THE 6.6 AI PRD

ACCEPTANCE CRITERIA Proposes valid slots: 80 of 100 golden cases pass all five rows , rubric v2 dated 12 Sep.	SUCCESS METRIC, PAIRED Sent with the proposed slot unedited, beside the weekly pass rate on a sample of real requests.	Every row grades a sentence . None grades an action . The day the assistant stops proposing a slot and starts booking it, a wrong answer becomes a wrong action and these five rows do not notice.
--	---	--

THE GAP THIS LEAVES — MODULE 8