

Build the leadership pack

Five moves, about four hours, one page at the end of it. Set a timer per move: going over is not thoroughness, it is the scope creep that kills self-paced work.

WHERE THIS GETS YOU • STEP 2 OF 3

1 — 2 — 3

When this is done you hold: a one-page decision memo, a quality number you produced yourself, a cost figure with token counts under it, a named gate, two honest unknowns, and five beats said out loud. **Banked: all 100 points. Carries into AC.3:** the page is what you present, and the two unknowns are what you learn to say first.

THE GATE • THREE THINGS, OR YOU ARE REPEATING THE MODULE RATHER THAN FINISHING IT

YOU CHOOSE

One real feature

Roadmap, vendor pitch, or a proposal you have never had evidence for. No candidate? Use the Rowan First Week Coach brief on the AC.1 page.

FROM M6 - M8

Spec, rubric, agent page

Adapted to this feature. Without them the quality block is thin, and that block is 25 points.

YOU BOOK IT

Two quiet hours, twice

Move 1 sets everything downstream. Do not start it in the twenty minutes before a meeting.

WRITE THIS FIRST, AND KEEP IT AT THE TOP OF THE FILE

I am asking _____ to decide _____ by _____.

Cannot fill all three slots? You have a status update. Fix that before move 2.

THE FIVE MOVES • ABOUT 4 HOURS • SET A TIMER PER MOVE

1 Assemble and decide 30 MIN • EARNS: THE DECISION, 20

Four artifacts side by side. Then three lines: **the decision** (a verb, a scope and a date — "approve a pilot with the 200 beta customers from 3 November", not "get alignment"), **the owner** (one person, by name, who says yes), and **the ask** (an engineer for two weeks, a budget line, a legal slot, a date).

PROMPT • PASTE YOUR SPEC FIRST

"Here is the spec for a feature I want to take to leadership. Draft three different decisions I could reasonably be asking for, from most cautious to most ambitious. For each, say what evidence it would require me to have and what it commits the company to. **Do not pick for me.**"

Done when: a stranger who reads only your first two lines knows what you want and by when.

2 The number that is against you 60 MIN • EARNS: QUALITY, 25

In this order. (1) Write the ship bar **before** you look at anything — "80 of 100 pass all five rows". A bar written afterwards is a rationalisation and readers can smell it. (2) Run your rubric over your golden set. No set? Hand-grade thirty real outputs this afternoon. (3) Get a number from **reality** — the same rubric over live traffic, or over outputs produced for people who are not you. No production number? Write that sentence on the page, with the date you will have one. **An absent number that is named is credible; an absent number that is hidden is failure mode three.** (4) Rank the failures by count, not by how annoying they are.

PROMPT • AFTER YOU HAVE RUN IT

"Here are my rubric rows and the results over [n] examples, including the failures. Group the failures, count each group, and tell me which group, if fixed, would move the overall score most. Then tell me **which of my rows could not have caught the failures I am seeing.**"

Done when: the least flattering number you have sits above the most flattering one.

3 The money 45 MIN • EARNS: UNIT ECONOMICS, 20

Token counts from a **real** request, not an estimate. Cost per request, per user per month, per month at your forecast volume — and say where the volume number came from; if it is a guess, use the word. Then the ten-times question (more users, longer inputs, more calls, a bigger model — which is most likely?), the lever you would pull first (cache, batch, route), and latency as the user feels it, not time to first token.

PROMPT • THEN CHECK THE ARITHMETIC YOURSELF

"Here is my per-request token count and the price per million tokens. Work out cost per request, per user per month at [n] uses, and the monthly total at [volume]. **Show the arithmetic so I can check it.** Then list what would make this ten times larger, in order of likelihood."

A wrong cost figure in a leadership memo is not survivable, and this is the easiest number in the pack to get wrong by a factor of a thousand.

Moves four and five, and the self-check

The risk block, the page itself, and the five beats out loud. Then score yourself against the card before anybody else does.

4 When it is wrong 45 MIN · EARNS: RISK AND CONTROL, 20

(1) **Name the worst plausible wrong action** as a sentence about a real person: "it sends an invite to a client for a time that does not work, and it cannot be recalled." Not "quality issues". (2) Place it on the Lesson 6.4 square: how bad, how permanent. (3) **The gate** — where a person stands, and why there. No gate? Say why, and what stands in its place: what the model may not say, what is computed rather than generated, what is checked before it goes out. (4) **What it may not touch** — the capabilities you withheld, as a list. This is the most convincing paragraph in most of these memos. (5) **What you tested** — ten requests, three runs each, and one attempt to make it misbehave.

PROMPT · THE ADVERSARIAL PASS

"Act as the person who would have to explain this feature's worst failure to a customer. Given this spec and these tools, list the ten worst things it could plausibly do, ordered by how hard each would be to undo. For each, say whether my design **prevents it, detects it, or neither.**"

"Neither" is the useful answer. Write those down as risks rather than deleting them.

5 The page, and the five minutes 60 MIN · EARNS: JUDGEMENT, 15

Write the page to the template: six blocks, one page, appendix underneath. Then the **unknowns** (two or three, each with what would settle it and how long that takes) and the **kill line** (what would make you stop this). Then the judgement paragraph. Then say the five beats **out loud, timed** — 60 / 90 / 60 / 60 / 30. Out loud, not in your head; you will cut two things the first time. Finish by running the ten questions from page 3 of the worked pack against your own page and marking the ones you cannot answer: those are not gaps in your preparation, they are the unknowns block.

SELF-CHECK · TICK HONESTLY, THEN SCORE THE CARD ON THE AC.1 PAGE

- | | |
|--|--|
| ☐ The first two lines name a decision, a date and an owner . | ☐ Cost per request has a token count under it, and I checked the arithmetic. |
| ☐ The ask is sized in people, money or time. | ☐ The ten-times risk and one lever are named. |
| ☐ The ship bar is dated before the number. | ☐ Latency is what a user waits for, not first token. |
| ☐ There is a number I produced by running something. | ☐ The worst wrong action is a sentence about a person. |
| ☐ The real-traffic number (or its named absence) is above the offline one. | ☐ There is a gate, or an argued reason there is none. |
| ☐ Failures are counted and ranked, not adjectives. | ☐ Withheld capabilities are listed, and I ran one misbehaviour test. |
| | ☐ Two unknowns, each with what would settle it, and a kill line. |
| | ☐ I have said it out loud, inside five minutes, without reading. |

THE JUDGEMENT LINE · 15 POINTS · FOUR MINUTES · WRITE IT NOW, NOT AT THE END

What the A I got wrong in this pack, and what I did about it:

The shape (yours will be specific): *the first cost model assumed one call per request and the real flow makes three, which tripled the figure · it proposed a rubric row for "helpfulness" and I cut it, because nobody has ever filed a bug saying an answer was insufficiently helpful · it wanted approval on every step and I moved the gate to the send, because approving nine steps is how people stop reading.*

This paragraph is the only part of the pack that proves a person was driving. It is also the row almost nobody fills in — in either capstone.